

sirilasso-ridgemodel

August 21, 2025

```
[1]: import pandas as pd
import numpy as np

#Import graphical plotting libraries
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

#Import Linear Regression Machine Learning Libraries
from sklearn import preprocessing
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score
```

```
[2]: data = pd.read_csv(r'C:\Users\ttwrd\Downloads\car-mpg.csv')
data.head()
```

```
[2]:
```

	mpg	cyl	disp	hp	wt	acc	yr	origin	car_type \
0	18.0	8	307.0	130	3504	12.0	70	1	0
1	15.0	8	350.0	165	3693	11.5	70	1	0
2	18.0	8	318.0	150	3436	11.0	70	1	0
3	16.0	8	304.0	150	3433	12.0	70	1	0
4	17.0	8	302.0	140	3449	10.5	70	1	0

```
car_name
0  chevrolet chevelle malibu
1      buick skylark 320
2    plymouth satellite
3      amc rebel sst
4      ford torino
```

```
[3]: data = data.drop(['car_name'], axis = 1)
```

```
[4]: data
```

```
[4]:      mpg  cyl  disp  hp   wt   acc  yr  origin  car_type
0    18.0   8  307.0  130  3504  12.0  70     1         0
1    15.0   8  350.0  165  3693  11.5  70     1         0
2    18.0   8  318.0  150  3436  11.0  70     1         0
3    16.0   8  304.0  150  3433  12.0  70     1         0
4    17.0   8  302.0  140  3449  10.5  70     1         0
..    ...  ...
393  27.0   4  140.0   86  2790  15.6  82     1         1
394  44.0   4   97.0   52  2130  24.6  82     2         1
395  32.0   4  135.0   84  2295  11.6  82     1         1
396  28.0   4  120.0   79  2625  18.6  82     1         1
397  31.0   4  119.0   82  2720  19.4  82     1         1
```

[398 rows x 9 columns]

```
[5]: data['origin'] = data['origin'].replace({1: 'america', 2: 'europe', 3: 'asia'})
```

```
[6]: data
```

```
[6]:      mpg  cyl  disp  hp   wt   acc  yr  origin  car_type
0    18.0   8  307.0  130  3504  12.0  70  america         0
1    15.0   8  350.0  165  3693  11.5  70  america         0
2    18.0   8  318.0  150  3436  11.0  70  america         0
3    16.0   8  304.0  150  3433  12.0  70  america         0
4    17.0   8  302.0  140  3449  10.5  70  america         0
..    ...  ...
393  27.0   4  140.0   86  2790  15.6  82  america         1
394  44.0   4   97.0   52  2130  24.6  82  europe          1
395  32.0   4  135.0   84  2295  11.6  82  america         1
396  28.0   4  120.0   79  2625  18.6  82  america         1
397  31.0   4  119.0   82  2720  19.4  82  america         1
```

[398 rows x 9 columns]

```
[7]: data = pd.get_dummies(data, columns = ['origin'])
```

```
[8]: data.head()
```

```
[8]:      mpg  cyl  disp  hp   wt   acc  yr  car_type  origin_america  \
0    18.0   8  307.0  130  3504  12.0  70         0             True
1    15.0   8  350.0  165  3693  11.5  70         0             True
2    18.0   8  318.0  150  3436  11.0  70         0             True
3    16.0   8  304.0  150  3433  12.0  70         0             True
4    17.0   8  302.0  140  3449  10.5  70         0             True

      origin_asia  origin_europe
0             False             False
```

```

1      False      False
2      False      False
3      False      False
4      False      False

```

```
[9]: data = data.replace('?', np.nan)
```

```
[10]: data
```

```
[10]:      mpg  cyl  disp  hp  wt  acc  yr  car_type  origin_america  \
0    18.0   8  307.0  130  3504  12.0  70         0             True
1    15.0   8  350.0  165  3693  11.5  70         0             True
2    18.0   8  318.0  150  3436  11.0  70         0             True
3    16.0   8  304.0  150  3433  12.0  70         0             True
4    17.0   8  302.0  140  3449  10.5  70         0             True
..    ...  ...  ...  ...  ...  ...  ...         ...             ...
393  27.0   4  140.0   86  2790  15.6  82         1             True
394  44.0   4   97.0   52  2130  24.6  82         1            False
395  32.0   4  135.0   84  2295  11.6  82         1             True
396  28.0   4  120.0   79  2625  18.6  82         1             True
397  31.0   4  119.0   82  2720  19.4  82         1             True

      origin_asia  origin_europe
0             False             False
1             False             False
2             False             False
3             False             False
4             False             False
..             ...             ...
393            False             False
394            False              True
395            False             False
396            False             False
397            False             False

```

```
[398 rows x 11 columns]
```

```
[26]: data=data.apply(pd.to_numeric,errors='ignore')
```

```

C:\Users\ttward\AppData\Local\Temp\ipykernel_67368\3768586041.py:1:
FutureWarning: errors='ignore' is deprecated and will raise in a future version.
Use to_numeric without passing `errors` and catch exceptions explicitly instead
data=data.apply(pd.to_numeric,errors='ignore')

```

```
[13]: data
```

```
[13]:      mpg  cyl  disp    hp  wt   acc  yr  car_type  origin_america  \
0    18.0   8  307.0  130.0 3504  12.0  70         0             True
1    15.0   8  350.0  165.0 3693  11.5  70         0             True
2    18.0   8  318.0  150.0 3436  11.0  70         0             True
3    16.0   8  304.0  150.0 3433  12.0  70         0             True
4    17.0   8  302.0  140.0 3449  10.5  70         0             True
..    ...   ...   ...   ...   ...   ...   ..         ...
393  27.0   4  140.0   86.0 2790  15.6  82         1             True
394  44.0   4   97.0   52.0 2130  24.6  82         1            False
395  32.0   4  135.0   84.0 2295  11.6  82         1             True
396  28.0   4  120.0   79.0 2625  18.6  82         1             True
397  31.0   4  119.0   82.0 2720  19.4  82         1             True
```

```
      origin_asia  origin_europe
0             False             False
1             False             False
2             False             False
3             False             False
4             False             False
..             ...             ...
393            False            False
394            False             True
395            False            False
396            False            False
397            False            False
```

[398 rows x 11 columns]

```
[27]: numeric_cols=data.select_dtypes(include=[np.number]).columns
data[numeric_cols] = data[numeric_cols].apply(lambda x: x.fillna(x.median()))
```

```
[28]: data.head()
```

```
[28]:      mpg  cyl  disp    hp  wt   acc  yr  car_type  origin_america  \
0    18.0   8  307.0  130.0 3504  12.0  70         0             True
1    15.0   8  350.0  165.0 3693  11.5  70         0             True
2    18.0   8  318.0  150.0 3436  11.0  70         0             True
3    16.0   8  304.0  150.0 3433  12.0  70         0             True
4    17.0   8  302.0  140.0 3449  10.5  70         0             True
```

```
      origin_asia  origin_europe
0             False             False
1             False             False
2             False             False
3             False             False
4             False             False
```

```
[29]: X = data.drop(['mpg'], axis = 1)
      y = data[['mpg']]
```

```
[30]: X_s = preprocessing.scale(X)
      X_s = pd.DataFrame(X_s, columns = X.columns)

      y_s = preprocessing.scale(y)
      y_s = pd.DataFrame(y_s, columns = y.columns) #
```

```
[31]: X_s
```

```
[31]:
```

	cyl	displacement	horsepower	weight	acceleration	year	car_type
0	1.498191	1.090604	0.673118	0.630870	-1.295498	-1.627426	-1.062235
1	1.498191	1.503514	1.589958	0.854333	-1.477038	-1.627426	-1.062235
2	1.498191	1.196232	1.197027	0.550470	-1.658577	-1.627426	-1.062235
3	1.498191	1.061796	1.197027	0.546923	-1.295498	-1.627426	-1.062235
4	1.498191	1.042591	0.935072	0.565841	-1.840117	-1.627426	-1.062235
..
393	-0.856321	-0.513026	-0.479482	-0.213324	0.011586	1.621983	0.941412
394	-0.856321	-0.925936	-1.370127	-0.993671	3.279296	1.621983	0.941412
395	-0.856321	-0.561039	-0.531873	-0.798585	-1.440730	1.621983	0.941412
396	-0.856321	-0.705077	-0.662850	-0.408411	1.100822	1.621983	0.941412
397	-0.856321	-0.714680	-0.584264	-0.296088	1.391285	1.621983	0.941412
..
393	0.773559	-0.497643	-0.461968				
394	0.773559	-0.497643	-0.461968				
395	0.773559	-0.497643	-0.461968				
396	0.773559	-0.497643	-0.461968				
397	0.773559	-0.497643	-0.461968				
..				
393	0.773559	-0.497643	-0.461968				
394	-1.292726	-0.497643	2.164651				
395	0.773559	-0.497643	-0.461968				
396	0.773559	-0.497643	-0.461968				
397	0.773559	-0.497643	-0.461968				

[398 rows x 10 columns]

```
[32]: y_s
```

```
[32]:
```

	mpg
0	-0.706439
1	-1.090751
2	-0.706439
3	-0.962647
4	-0.834543

```

..      ...
393  0.446497
394  2.624265
395  1.087017
396  0.574601
397  0.958913

```

```
[398 rows x 1 columns]
```

```
[34]: X_train, X_test, y_train, y_test = train_test_split(X_s, y_s, test_size = 0.30,
↳ random_state = 1)
X_train.shape
```

```
[34]: (278, 10)
```

```
[35]: regression_model = LinearRegression()
regression_model.fit(X_train, y_train)

for idx, col_name in enumerate(X_train.columns):
    print('The coefficient for {} is {}'.format(col_name, regression_model.
↳ coef_[0][idx]))

intercept = regression_model.intercept_[0]
print('The intercept is {}'.format(intercept))
```

```

The coefficient for cyl is 0.3210223856916108
The coefficient for disp is 0.3248343091848394
The coefficient for hp is -0.2291695005943759
The coefficient for wt is -0.7112101905072299
The coefficient for acc is 0.014713682764191435
The coefficient for yr is 0.3755811949510741
The coefficient for car_type is 0.38147694842331
The coefficient for origin_america is -0.0747224754758417
The coefficient for origin_asia is 0.04451525203567813
The coefficient for origin_europe is 0.04834854953945371
The intercept is 0.019284116103639715

```

```
[37]: ridge_model = Ridge(alpha = 0.4)
ridge_model.fit(X_train, y_train)

print('Ridge model coef: {}'.format(ridge_model.coef_))
```

```

Ridge model coef: [ 0.31495967  0.30948411 -0.22861679 -0.69782283  0.01239531
0.37411266
 0.37586629 -0.07408168  0.04437854  0.0476772 ]

```

```
[40]: lasso_model = Lasso(alpha = 0.1)
lasso_model.fit(X_train, y_train)

print('Lasso model coef: {}'.format(lasso_model.coef_))
```

```
Lasso model coef: [-0.          -0.          -0.01690287 -0.51890013  0.
0.28138241
0.1278489  -0.01642647  0.          0.          ]
```

```
[39]: print(regression_model.score(X_train, y_train))
print(regression_model.score(X_test, y_test))

print('*****')
#Ridge
print(ridge_model.score(X_train, y_train))
print(ridge_model.score(X_test, y_test))

print('*****')
#Lasso
print(lasso_model.score(X_train, y_train))
print(lasso_model.score(X_test, y_test))
```

```
0.8343770256960538
0.8513421387780067
*****
0.8343502868181134
0.8520594956782537
*****
0.7938010766228453
0.8375229615977084
```

```
[ ]:
```